

# Computational Communication Science 2

## Week 7 - Lecture

### »Validation (in Supervised Machine Learning)«

---

a.m.moller@uva.nl

May 15, 2023

Digital Society Minor, University of Amsterdam

# Today

Recap

Validating models

Validation metrics

About input for SML

## Recap

---

# Recap

## Last week, we discussed:

- Supervised Machine Learning (SML)
- The principles behind SML
- The steps of SML
- Some commonly used ML models

## At home, you:

- Got some hands-on experience SML (week 6 exercises)

# Recap

## Today, we:

- Review your first SML experience
- Take a deep dive into validating SML-models

# Recap

## Last week, you practiced with code that:

- Read in some data (Q1)
- Split the data into a train and a test set (Q2)
- Set up a Count vectorizer (Q3)
- Trained a Naïve Bayes model with the count vectorizer (Q4)
- Requested some metrics for validation (Q5)

# Recap Q1

```
1 import csv
2 from collections import Counter
3 import matplotlib.pyplot as plt
4
5 file = "hatespeech_text_label_vote_RESTRICTED_100K.csv"
6 tweets = []
7 labels = []
8
9 with open(file) as fi:
10     data = csv.reader(fi, delimiter='\t')
11     for row in data:
12         tweets.append(row[0])
13         labels.append(row[1])
14
15 print(len(tweets) == len(labels)) # there should be just as many tweets
    as there are labels
16
17 Counter(labels)
18 plt.bar(Counter(labels).keys(), Counter(labels).values())
```

## Recap Q2

Split the dataset:

```
1 from sklearn.model_selection import train_test_split
2
3 tweets_train, tweets_test, y_train, y_test = train_test_split(tweets,
    labels, test_size=0.2, random_state=42)
```



# Recap Q3

What happens here?

```
1 from sklearn.feature_extraction.text import (CountVectorizer)
2
3 countvectorizer = CountVectorizer(stop_words="english")
4 X_train = countvectorizer.fit_transform(tweets_train)
5 X_test = countvectorizer.transform(tweets_test)
```

# Recap Q4

The actual SML part (yes, truly, it is three lines of code!):

```
1 nb = MultinomialNB()  
2 nb.fit(X_train, y_train)  
3 y_pred = nb.predict(X_test)
```

# Recap Q5

You can check what was created:

```
1 nb = MultinomialNB()
2 nb.fit(X_train, y_train)
3 y_pred = nb.predict(X_test)
4
5 print(y_pred[:10])
```

```
1 ['normal' 'normal' 'normal' 'normal' 'spam' 'normal' 'normal' '
   normal' 'abusive' 'normal']
```

## Recap Q5

Classification report:

```
1 from sklearn.metrics import classification_report
2
3 print(classification_report(y_test, y_pred))
```

## Up next

Classification report: validate your classifier.

More about this today!

## Validating models

---

# Validation

Validation: When we assess the performance of a classifier.

Or when we try to answer the question: "How well does the classifier work?"

# Validation

Validation: When we assess the performance of a classifier.

Or when we try to answer the question: "How well does the classifier work?"



# Validation

What criteria should we use to decide on this?

Entirely context specific!

# Validation

What criteria should we use to decide on this?

Entirely context specific!

# Validation

## Compare different goals for using SML:

- To automatically decide what Instagram users should see an advertisement
- To automatically remove spam from Twitter feed

Would you use the same criterion in both cases to determine how well a classifier works? Why (not)?

# Validation

There are various evaluation metrics available for machine learning.  
In scikit-learn, they are presented by ways of a classification report!

# Zooming out

## So far, we:

- Reviewed the exercise and the basic steps of SML
- Talked about what validation is

## Next, we will talk about:

- Some commonly used validation metrics
- Input for SML
- Finding the best classifier

## Validation metrics

---

# Precision

Precision quantifies the number of positive class predictions that actually belong to the positive cases.

OR: How much of what we found is actually correct?

# Precision

Precision quantifies the number of positive class predictions that actually belong to the positive cases.

OR: How much of what we found is actually correct?



# Precision

Precision quantifies the number of positive class predictions that actually belong to the positive cases.

OR: How much of what we found is actually correct?

## Compare different goals for using SML in exercise 6:

- To automatically decide what Instagram users should see an advertisement
- To automatically remove spam from Twitter feed

# Recall

Recall quantifies the number of positive class prediction made out of all positive examples in the dataset.

OR: How many of the cases that we wanted to find did we actually find?

# Recall

Recall quantifies the number of positive class prediction made out of all positive examples in the dataset.

OR: How many of the cases that we wanted to find did we actually find?

# Recall

Recall quantifies the number of positive class prediction made out of all positive examples in the dataset.

OR: How many of the cases that we wanted to find did we actually find?

## Compare different goals for using SML in exercise 6:

- To automatically decide what Instagram users should see an advertisement
- To automatically remove spam from Twitter feed

## Precision and Recall

		Predicted class	
		0	1
Actual class	0	True Neg	False Pos
	1	False Neg	True Pos

# Precision and Recall

		Predicted class	
		0	1
Actual class	0	180 (TN)	50 (FP)
	1	20 (FN)	150 (TP)

Precision is calculated as:  $\frac{TP}{TP+FP}$

In this example  $\frac{150}{150+50}$  which is 0.75

Recall is calculated as  $\frac{TP}{TP+FN}$

In this example  $\frac{150}{150+20}$  which is 0.88

# What does this look like in code?

Let's ask for a confusion matrix:

```
1 from sklearn.metrics import confusion_matrix
2
3 y_test = [0, 1, 1, 1, 0]
4 y_pred = [0, 0, 1, 1, 1]
5
6 print(confusion_matrix(y_test, y_pred))
```

```
1 [[1 1]
2  [1 2]]
```

# The classification report

Let's get some metrics for validation:

```
1 from sklearn.metrics import classification_report
2 print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
abusive	0.81	0.88	0.85	5369
hateful	0.83	0.05	0.10	966
normal	0.78	0.93	0.85	10848
spam	0.67	0.30	0.41	2817
accuracy			0.78	20000
macro avg	0.77	0.54	0.55	20000
weighted avg	0.78	0.78	0.75	20000



# $F_1$ -score

But wait...

## Compare different goals for using SML:

- To automatically decide what Instagram users should see an advertisement
- To automatically remove spam from Twitter feed

Such information was not available in the exercise for week 6!

# $F_1$ -score

$F_1$ -score: The harmonic mean of precision and recall. (Weighted average of precision and recall)

$$F_1\text{-score} = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

# Accuracy

	precision	recall	f1-score	support
abusive	0.81	0.88	0.85	5369
hateful	0.83	0.05	0.10	966
normal	0.78	0.93	0.85	10848
spam	0.67	0.30	0.41	2817
accuracy			0.78	20000
macro avg	0.77	0.54	0.55	20000
weighted avg	0.78	0.78	0.75	20000

# Accuracy


Accuracy: In which percentage of all cases was our classifier right?

# Accuracy


Class distribution: The number of examples that belong to each class.

Imbalanced classification: A predictive modeling problem where the distribution of examples across the classes within a training dataset is not equal.

# Accuracy



# Accuracy



Majority class (red dots) vs. minority class (grey dots)

# Accuracy

Always check how your cases are distributed across the labels.



## Validation: Best practices

Decide what metric is best to use beforehand.

### Consider:

- Is class imbalance an issue?
- What will the classifier be used for?

The latter can bring you back to the question: To SML or not to SML?

# Validation: Best practices

## SML suitability depends on:

- How hard/easy it is to translate the decision proces for classification into straight-forward rules
- How much data there are to classify
- How much room there is for errors

Consider these matters as you decide (a) whether or not to use SML, and (b) what performance metrics to use.

## Validation: Best practices

In last week's lecture, we saw that you can train many different classifiers.

### Amongst other, classifiers can differ based on:

- The vectorizer that is used on the data (i.e., count vectorizer or tf-idf vectorizer)
- The underlying model (e.g., Naïve Bayes, Logistic Regression, Decision Trees, SVM, etc.)

How do you know what classifier is best beforehand?

## Validation: Best practices

In last week's lecture, we saw that you can train many different classifiers.

### Amongst other, classifiers can differ based on:

- The vectorizer that is used on the data (i.e., count vectorizer or tf-idf vectorizer)
- The underlying model (e.g., Naïve Bayes, Logistic Regression, Decision Trees, SVM, etc.)

How do you know what classifier is best beforehand?

## Selecting a classifier

You don't!

Typically, various classifiers are trained and their performance is compared.

The best performing classifier is then selected and used to annotate more/new data.

## Selecting a classifier

You don't!

Typically, various classifiers are trained and their performance is compared.

The best performing classifier is then selected and used to annotate more/new data.

## Selecting a classifier

You don't!

Typically, various classifiers are trained and their performance is compared.

The best performing classifier is then selected and used to annotate more/new data.

# Zooming out

## So far, we:

- Reviewed the exercise and the basic steps of SML
- Talked about what validation is
- Discussed some commonly used validation metrics

## Next, we will talk about:

- Input for SML



## About input for SML

---

# Bad input

## Poor quality input can cause:

- Your classifier not being able to identify specific categories
- Your classifier to perform badly overall
- Your classifier to pick up implicit bias

# Bad input

What makes input (training data) bad input?

# Bad input

What can be done to solve these issues?

## Poor quality input can cause:

- Your classifier not being able to identify specific categories
- Your classifier to perform badly overall
- Your classifier to pick up implicit bias

# Creating good input

## You can create good input:

- Avoid class imbalance: increase your sample (add rare cases), or random selection
- Make sure the hand coding is reliable (check Krippendorff's Alpha, add coder training)
- Make sure to develop a clear and objective codebook

# The limits of your input

User case\*:

A classifier was trained to distinguish between YouTube comments that address the video that they accompany, and YouTube comments that do not. The classifier was trained on a training set consisting of > 16.000 comments written in response to music videos. The classifier performed well, with high scores on all evaluation metrics. The researchers now want to use their classifier to automatically label a new set of 10.000 YouTube comments written in response to vlogs. Can they?

\*Based on: Möller et al., in press

# The limits of your input

User case\*:

A classifier was trained to distinguish between YouTube comments that address the video that they accompany, and YouTube comments that do not. The classifier was trained on a training set consisting of  $> 16.000$  comments written in response to music videos. The classifier performed well, with high scores on all evaluation metrics. The researchers now want to use their classifier to automatically label a new set of 10.000 YouTube comments written in response to vlogs. Can they?

\*Based on: Möller et al., in press

# The limits of your input

User case\*:

A classifier was trained to distinguish between YouTube comments that address the video that they accompany, and YouTube comments that do not. The classifier was trained on a training set consisting of  $> 16.000$  comments written in response to music videos. The classifier performed well, with high scores on all evaluation metrics. The researchers now want to use their classifier to automatically label a new set of 10.000 YouTube comments written in response to vlogs. Can they?

\*Based on: Möller et al., in press



# The limits of your input

User case\*:

A classifier was trained to distinguish between YouTube comments that address the video that they accompany, and YouTube comments that do not. The classifier was trained on a training set consisting of  $> 16.000$  comments written in response to music videos. The classifier performed well, with high scores on all evaluation metrics. The researchers now want to use their classifier to automatically label a new set of 10.000 YouTube comments written in response to vlogs. *Can they?*

\*Based on: Möller et al., in press

## The limits of your input

User case\*:

A classifier was trained to distinguish between YouTube comments that address the video that they accompany, and YouTube comments that do not. The classifier was trained on a training set consisting of  $> 16.000$  comments written in response to music videos. The classifier performed well, with high scores on all evaluation metrics. The researchers now want to use their classifier to automatically label a new set of 10.000 YouTube comments written in response to vlogs. Can they?

\*Based on: Möller et al., in press

## The limits of your input

Classifiers can only be used to automatically annotate data from the same population as the data that they were trained on.

# Zooming out

## Today, we:

- Reviewed the exercise and the basic steps of SML
- Talked about what validation is
- Discussed some commonly used validation metrics
- Talked about input for SML

# Zooming out

## Tomorrow and this week, you will:

- Set up multiple different classifiers
- Validate those classifiers
- Select the best performing classifier

Work on the tutorial exercises for this week.